

Schedule

- 4/5 (Tue) Lecture – 3D rotation (Chap 6~8) / Interpolation (Chap 9)
- 4/6 (Wed) Open Lab
- 4/7 (Thur) Lecture – Projection& Depth (Chap10,11)

- 4/12 (Tue) TA's Special Session for OpenGL (RE: **HW#3 out!**)
- 4/13 <Election Day> No lab session ----- **HW#2 DUE**
- 4/14 (Thur) Lecture – From Vertex to Pixel (Chap.12)

- 4/19 (Tue) Lecture – Modeling (Chap.22) or Varying variable (Chap 13)
- 4/20~26 (Midterm Week)
- 4/26 (Tuesday) 4~7 PM **Midterm Exam @ E3-1 #1501**

- 4/27 (Wed) Open Lab
- 4/28 (Thur) Lecture – Lighting (Chap. 13) ----- **HW#3 DUE**

Questions from last lecture

$$R_\alpha := (R_1 R_0^{-1})^\alpha R_0$$

- About cn

$$\left(cn \left[\begin{bmatrix} \cos\left(\frac{\theta_1}{2}\right) \\ \sin\left(\frac{\theta_1}{2}\right) \hat{\mathbf{k}}_1 \end{bmatrix} \begin{bmatrix} \cos\left(\frac{\theta_0}{2}\right) \\ \sin\left(\frac{\theta_0}{2}\right) \hat{\mathbf{k}}_0 \end{bmatrix}^{-1} \right]^\alpha \begin{bmatrix} \cos\left(\frac{\theta_0}{2}\right) \\ \sin\left(\frac{\theta_0}{2}\right) \hat{\mathbf{k}}_0 \end{bmatrix} \right)$$
- Note that a rotation of $-\theta$ degrees about the axis $-\hat{\mathbf{k}}$ gives us the same quaternion.
- A rotation of $\theta + 4\pi$ degrees about an axis $\hat{\mathbf{k}}$ also gives us the **same** quaternion
- A rotation of $\theta + 2\pi$ degrees about an axis $\hat{\mathbf{k}}$, which in fact is the same rotation, gives us the **negated** quaternion
- So antipodes represent the same rotation transformation

Questions from last lecture


$$R_\alpha := (R_1 R_0^{-1})^\alpha R_0$$

- If the **transition quaternion** $\begin{bmatrix} \cos(\frac{\theta}{2}) \\ \sin(\frac{\theta}{2}) \hat{\mathbf{k}} \end{bmatrix}$ presents a θ of more than π (180) degrees, in particular, if $\cos(\frac{\theta}{2}) < 0$ then $\theta \in [\pi \dots 2\pi]$, $\alpha\theta$ would go more than 180 degrees which we don't want during interpolation
- In this case, suppose we had swapped to the antipode before calling power. Then $\cos(\frac{\theta}{2}) > 0$, we get $\theta/2 \in [-\pi/2 \dots \pi/2]$. And thus $\theta \in [-\pi \dots \pi]$
- So when we interpolate, before calling the power operator, we first check the sign of the first coordinate, and conditionally negate the quaternion.
- We call this the **conditional negation operator** cn

$$\left(cn \left[\begin{bmatrix} \cos(\frac{\theta_1}{2}) \\ \sin(\frac{\theta_1}{2}) \hat{\mathbf{k}}_1 \end{bmatrix} \right] \left[\begin{bmatrix} \cos(\frac{\theta_0}{2}) \\ \sin(\frac{\theta_0}{2}) \hat{\mathbf{k}}_0 \end{bmatrix} \right]^{-1} \right)^\alpha \left[\begin{bmatrix} \cos(\frac{\theta_0}{2}) \\ \sin(\frac{\theta_0}{2}) \hat{\mathbf{k}}_0 \end{bmatrix} \right]$$

jinah@cs.kaist.ac.kr

CS380



Gimbal Lock Problem

Extracting Euler angles

$$F = \begin{pmatrix} f_{00} & f_{01} & f_{02} \\ f_{10} & f_{11} & f_{12} \\ f_{20} & f_{21} & f_{22} \end{pmatrix} = R_z(r)R_x(p)R_y(h) = E(h,p,r)$$

With $\cos(a) = C_a$, $\sin(a) = S_a$

$$= \begin{pmatrix} C_r C_h - S_r S_p S_h & -S_r C_p & C_r S_h + S_r S_p C_h \\ S_r C_h + C_r S_p S_h & C_r C_p & S_r S_h - C_r S_p C_h \\ -C_p S_h & S_p & C_p C_h \end{pmatrix}$$

$$\frac{f_{01}}{f_{11}} = \frac{-\sin r}{\cos r} = -\tan r$$

$$\frac{f_{20}}{f_{22}} = \frac{-\sin h}{\cosh} = -\tan h$$

$$f_{21} = \sin p$$


$$r = \text{atan2}(-f_{01}, f_{11})$$

$$h = \text{atan2}(-f_{20}, f_{22})$$


$$p = \arcsin(f_{21})$$

Problem: what if $\cos p = 0$?

Some references on quaternion (paper) and gimbal lock (demo) are posted on KLMS



O' ?



$$\vec{o}^t = \vec{w}^t O$$

$$\vec{o}'^t = \vec{o}^t O'$$

$$\vec{a}^t = \vec{o}^t A$$

$$\vec{b}^t = \vec{a}^t B$$

$$\vec{b}'^t = \vec{b}^t B'$$

$$\vec{c}^t = \vec{b}^t C$$

$$\vec{d}^t = \vec{c}^t D$$

$$\vec{d}'^t = \vec{d}^t D'$$

$$\vec{f}^t = \vec{o}^t F$$

Projection & Depth

Chapter 10 & 11




Slide from Prof. MH Kim

Camera transforms

- Until now we have considered all of our geometry in a 3D space
- Ultimately everything ended up in eye coordinates with coordinates $[x_e, y_e, z_e, 1]^t$
- We said that the camera is placed at the origin of the eye frame \vec{e}^t , and that it is looking down the eye's negative z-axis.
- This *somehow* produces a 2D image.
- We had a *magic matrix* which created `gl_Position`
- Now we will study this step


Pinhole camera model



- As light travels towards the film plane, most is blocked by an opaque surface placed at the $z_e = 0$ plane.
- But we place a very small hole in the center of the surface, at the point with eye coordinates $[0,0,0,1]^t$

9


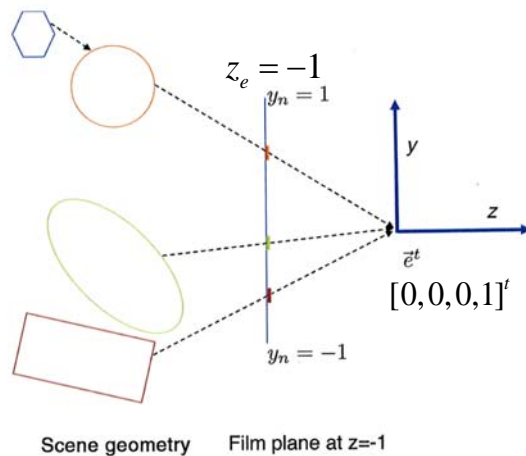
Pinhole camera model



- Only rays of light that pass through this point reach the film plane and have their intensity recorded on film. The image is recorded at a film plane placed at, say, $z_e = 1$

10



Pinhole camera model



- A physical camera needs a finite aperture and a lens, but we will ignore this.
- To avoid the image flip, we can mathematically model this with the film plane in front of the pinhole, say at the $z_e = -1$

11


Pinhole camera model



- If we hold up the photograph at the $z_e = -1$ plane, and observe it with our own eye, placed at the origin, it will look to us just like the origin scene would have.

12


Basic mathematical model



- Let us use **normalized** coordinates $[x_n, y_n]^t$ to specify points on our film plane.
 - For now, let them match **eye coordinates** on **this film plane**.
- Where does the ray from \tilde{p} to the origin hits the film plane?

13


Basic mathematical model



- All points on the ray hit the same pixel.
- All points on the ray are all scales
- So points on ray are: $[x_e, y_e, z_e]^t = \alpha[x_n, y_n, -1]^t$

14

Basic mathematical model




□ So $[x_e, y_e, z_e]^t = -z_e [x_n, y_n, -1]^t$

□ So
$$x_n = -\frac{x_e}{z_e}, \quad y_n = -\frac{y_e}{z_e}$$

15

Projection matrix




□ We can model this expression as a matrix operation.

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} = \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ - \\ w_c \end{bmatrix}$$

16


In matrix form



- The raw output of the matrix multiply, $[x_c, y_c, -, w_c]^t$ are called the **clip coordinates** of \tilde{p} .
- $w_n = w_c$ is a new variable called the **w-coordinate**.
 - In such clip coordinates, the fourth entry of the coordinate 4-vector is not necessarily a zero or a one.

17

Divide by w




- We say that $x_n w_n = x_c$ and $y_n w_n = y_c$. If we want to extract x_n alone, we must perform the division $x_n = \frac{x_n w_n}{w_n}$
- This recovers our camera model

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} = \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ - \\ w_c \end{bmatrix}$$

18


Divide by w



- Our output coordinates, with subscripts 'n', are called *normalized device coordinates (NDC)* because they address points on the image in abstract units without specific reference to numbers of pixels.

19


Divide by w



- We keep all of the image data in the *canonical square*, $-1 \leq x_n \leq +1, -1 \leq y_n \leq +1$, and ultimately map this onto a window on the screen.
 - Data outside of this square does not be recorded or displayed.
 - This is exactly the model we used to describe 2D OpenGL

20


Scales



- By changing the entries in the projection matrix, we can slightly alter geometry of the camera transformation.
- We could push the film plane out to $z_e = n$, where n is some negative number (zoom lens)

21


Scales



- So points on ray are: $[x_e, y_e, z_e]^t = \alpha[x_n, y_n, z_n]^t$
- So $[x_e, y_e, z_e]^t = \frac{z_e}{n}[x_n, y_n, z_n]^t$
- So $x_n = \frac{x_e n}{z_e}, y_n = \frac{y_e n}{z_e}$

22

In matrix form




□ In matrix form, this becomes:

(supposing n is some negative number)

$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} -n & 0 & 0 & 0 \\ 0 & -n & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

23

In matrix form




□ Note this matrix is the same as

$$\begin{bmatrix} -n & 0 & 0 & 0 \\ 0 & -n & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

24


In matrix form



- This has the same effect as starting with our original camera, scaling by $-n$, and cropping to the canonical square.

25

fovY




- Scale can be determined by *vertical angular field of view* of the desired camera.
- If we want our camera to have a field of view of θ degrees, then we can set $-n = \frac{1}{\tan\left(\frac{\theta}{2}\right)}$ giving us

26

fovY

- Verify that any point whose ray from the origin forms a vertical angle of $\theta/2$ with the negative z axis maps to the boundary of the canonical square



- The point with eye coordinates: $[0, \tan(\frac{\theta}{2}), -1, 1]^t$ maps to

normalized device coordinates

$[0, 1]^t$

$$\begin{bmatrix} \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

27

Dealing with aspect ratio

- Suppose the window is wider than its height. In our camera transform, we need to squish things horizontally so a wider horizontal field of view fits into our retained canonical square.
- When the data is later mapped to the window, it will be stretched out correspondingly and will not appear distorted.
- Define a , the *aspect ratio* of a window, to be its width divided by its height (measured say in pixels).

$$a = \frac{(\text{width px})}{(\text{height px})}$$

28

Dealing with aspect ratio

- We can then set our projection matrix to be:

$$\begin{bmatrix} \frac{1}{\alpha \tan\left(\frac{\theta}{2}\right)} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan\left(\frac{\theta}{2}\right)} & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

- So when the window is wide, we will keep more horizontal FOV, and when the window is tall, we will keep less horizontal FOV.


29

FOV issues

- To be a “window” onto the world, the FOV should match the angular extents of the window in the viewers field.
- This might give a too limited view onto the world.
- So we can increase it to see more.
- But this might give a somewhat unnatural look.

30


Shifts



- Sometimes, we wish to crop the image non-centrally.
- This can be modeled as translating the **normalized device coordinates** (NDC)'s and then cropping centrally.

31

Shifts




$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & c_x \\ 0 & 1 & 0 & c_y \\ - & - & - & - \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

$$= \begin{bmatrix} 1 & 0 & -c_x & 0 \\ 0 & 1 & -c_y & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

32

Shifts



- Useful for tiled displays, stereo viewing, and certain kinds of images.



33

Frustum


- Shifts are often specified by first specifying a near plane.

$$z_e = n$$

- On this plane, a rectangle is specified with the eye coordinates of an axis aligned rectangle. (for non-distorted output, the aspect ratio of this rectangle should match that of the final window.)

- Using l, r, t, b .

$$\begin{bmatrix} -\frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & -\frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$




34

Context

- Projection could be applied to every point in the scene.
- In CG, we will apply it to the vertices to position a triangle on the screen.
- The rest of the triangle will then get filled in on the screen as we shall see.

35


Summary: Pinhole camera model



- Only rays of light that pass through this point reach the film plane and have their intensity recorded on film. The image is recorded at a film plane placed at, say, $z_e = 1$

36

Summary: Normalized device coordinates




□ Canonical square space:

$$x_n = -\frac{x_e}{z_e}, \quad y_n = -\frac{y_e}{z_e} \quad \xrightarrow{w_n} \quad \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ - \\ w_c \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

37

Summary: Scale factor n




□ Controlling aspect ratio of film space


$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} -n & 0 & 0 & 0 \\ 0 & -n & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix} \quad \rightarrow \quad \begin{bmatrix} x_n w_n \\ y_n w_n \\ - \\ w_n \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

Summary:

Frustum: Eye coord. → NDC




$$\begin{bmatrix} \frac{1}{\alpha \tan\left(\frac{\theta}{2}\right)} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan\left(\frac{\theta}{2}\right)} & 0 & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$



$$\begin{bmatrix} -\frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & -\frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ - & - & - & - \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

39


Visibility



- ❑ In the real world, opaque objects block light.
- ❑ We need to model this computationally.
- ❑ One idea is to render back to front and use overwriting
 - This will have problem with visibility cycles.

40

Visibility



- We could explicitly store everything hit along a ray and then compute the closest.
 - Make sense in a **ray tracing** setting, where we are working **one pixel per ray at time**, but not for OpenGL, where we are working **one triangle at a time**.

41

Z-buffer

- We will use z-buffer
- Triangle are drawn in **any** order
- Each pixel in frame buffer stores 'depth' value of closest geometry observed so far.
- When a new triangle tries to set the color of a pixel, we first compare its depth to the value stored in the z-buffer.
- Only if the observed point in this triangle is closer, we overwrite the color and depth values of this pixel.
- This is done **per-pixel**, so there is no cycle problems.
- There are optimizations, where z-testing is done, before the fragment shading is done.

42

Other uses of visibility calculations

- Visibility to a light source is useful for shadows.
 - We will talk about shadow mapping later.
 - We will do shadow calculations in a ray tracer.
- Visibility computation can also be used to speed up the rendering process.
 - If we know that some object is occluded from the camera, then we don't have to render the object in the first place.
 - We can use a conservative test.

43

Basic mathematical model

- For every point, we define its $[x_n, y_n, z_n]^t$ coordinates, using the following matrix expression:

$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ z_n w_n \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix} = \begin{bmatrix} s_x & 0 & -c_x & 0 \\ 0 & s_y & -c_y & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

- We now also have the value $z_n = \frac{-1}{z_e}$
- Our plan is to use this z_n value to do depth comparison in our z-buffer.

44


Correct ordering

- Given two points \tilde{p}^1 and \tilde{p}^2 with eye coordinates $[x_e^1, y_e^1, z_e^1, 1]^t$ and $[x_e^2, y_e^2, z_e^2, 1]^t$.
- Suppose that they both are in front of the eye, i.e., $z_e^1 < 0$ and $z_e^2 < 0$.
- And suppose that \tilde{p}^1 is closer to the eye than \tilde{p}^2 , that is $z_e^2 < z_e^1$.
- Then

$$-\frac{1}{z_e^2} < -\frac{1}{z_e^1},$$

meaning

$$z_e^2 < z_e^1$$




Projective transform

- We can now think of the process of taking points (given by **eye coordinates**) to points (given by **normalized device coordinates**) as an honest-to-goodness 3D geometric transformation.
- This kind of transformation is generally neither linear nor affine, but is something called a **3D projective transformation**.
- Projective transformation preserve **co-linearity** and **co-planarity** of points.

Co-linearity of points

- If three or more points are on a single line, the transformed points will also be on some single line.





- Three points $\mathbf{x}_i = [x_i, y_i, z_i, 1]$ for $i = 1, 2, 3$

$$x_2 - x_1 : y_2 - y_1 : z_2 - z_1 = x_3 - x_1 : y_3 - y_1 : z_3 - z_1$$

$$\|(\mathbf{p}_2 - \mathbf{p}_1) \times (\mathbf{p}_3 - \mathbf{p}_1)\| = 0$$

47

Co-planarity of points




- Note that distances are not preserved by a projective transform.
- Evenly spaced pixel on the film do not correspond to evenly spaced points on the geometry in eye space.
- Meanwhile, such *evenly spaced pixels* correspond with evenly spaced points in *normalized device coordinates*.

48

z_n interpolation is right

- Preservation of coplanarity: for points on a fixed triangle, we will have $z_n = ax_n + by_n + c$, for some fixed a, b and c
- Thus, the correct z_n value for a point can be computed using linear interpolation over the 2D image domain as long as we know its value at the three vertices of the triangle.



49

Solution: near/far

- Solution: replacing the third row of the matrix with more general row $\begin{bmatrix} 0 & 0 & \alpha & \beta \end{bmatrix}$

$$\rightarrow \begin{bmatrix} x_n w_n \\ y_n w_n \\ z_n w_n \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix} = \begin{bmatrix} s_x & 0 & -c_x & 0 \\ 0 & s_y & -c_y & 0 \\ 0 & 0 & \alpha & \beta \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$

- It is easy to verify that if the value α and β are both positive, then the z-ordering of points (assuming they all have negative z_e values) is preserved under the projective transform.

50


Solution: near/far

- To set α and β , we first select depth value n and f called the *near* and *far* values (both negative), such that our main region of interest in the scene is sandwiched between $z_e = n$ and $z_e = f$

- Given these selections we set

$$\alpha = \frac{f + n}{f - n}$$


$$\beta = -\frac{2fn}{f - n}$$



51


Solution: near/far

- We can verify now that any point with $z_e = f$ maps to a point with $z_n = -1$ and that a point with $z_e = n$ maps to a point with $z_n = 1$
- Any geometry not in this [near...far] range is clipped away by OpenGL and ignored.



52

Proj. Trans.: Eye coor. → NDC



$$\begin{bmatrix} \frac{1}{\alpha \tan\left(\frac{\theta}{2}\right)} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan\left(\frac{\theta}{2}\right)} & 0 & 0 \\ 0 & 0 & \frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} -\frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & -\frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & \frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$


53

Codes

- ❑ In OpenGL, use the z-buffer is turned on with a call to `glEnable(GL_DEPTH_TEST)`.
- ❑ We also need a call to `glDepthFunc(GL_GREATER)`, since we are using a right handed coordinate system where 'more-negative' is 'farther from the eye'.
- ❑ In real life, you may see other conventions (for how to interpret n and f , some of the signs of the matrix, and the handedness of the ultimate z-test).

54

Summary: Visibility



- We could explicitly store everything hit along a ray and then compute the closest.
 - Make sense in a **ray tracing** setting, where we are working **one pixel per ray at a time**, but not for OpenGL, where we are working **one triangle at a time**.

55

Summary: Z-buffer

- We will use z-buffer
- Triangle are drawn in any order
- **Each pixel** in frame buffer stores **'depth' value of closest geometry** observed so far.
- When a new triangle tries to set the color of a pixel, we first compare its depth to the value stored in the z-buffer.
- Only if the observed point in this triangle is closer, we overwrite the color and depth values of this pixel.

56


Summary:

Proj. Trans.: Eye coor. → NDC

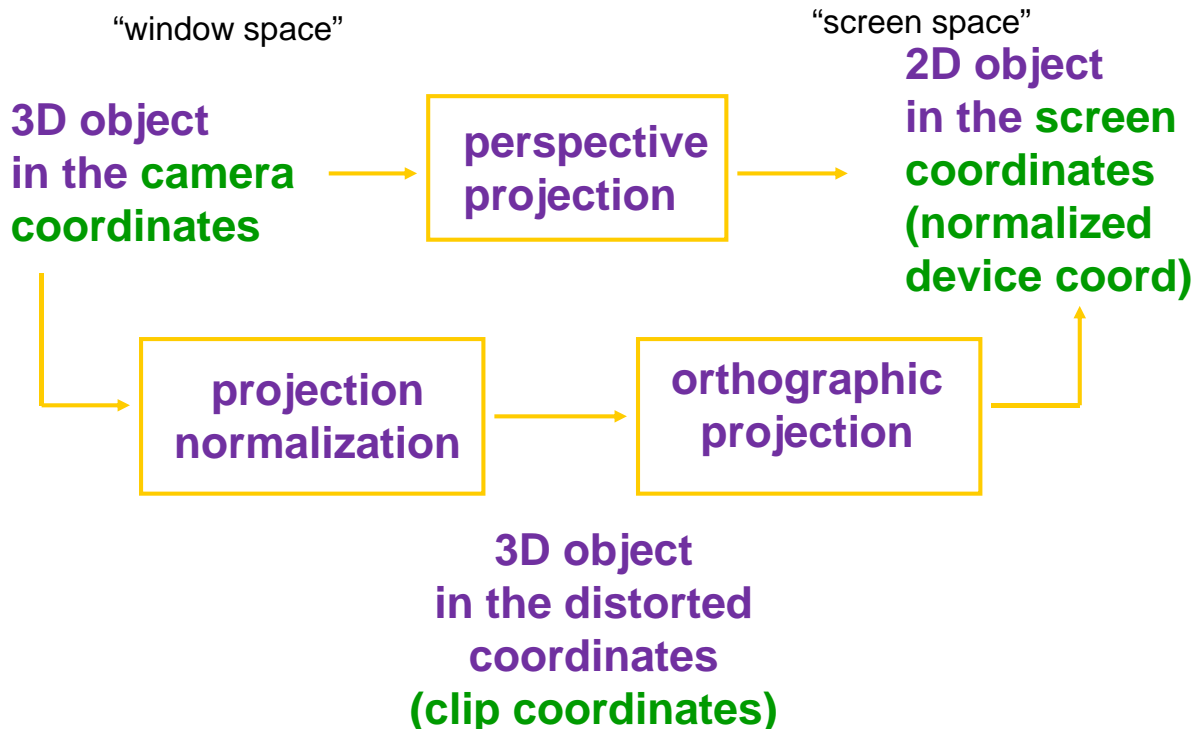
□ Camera projection transformation

$$\begin{bmatrix} x_n w_n \\ y_n w_n \\ z_n w_n \\ w_n \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix} = \begin{bmatrix} -\frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & -\frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & \frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x_e \\ y_e \\ z_e \\ 1 \end{bmatrix}$$


$$M = \begin{bmatrix} \frac{1}{\alpha \tan(\frac{\theta}{2})} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\frac{\theta}{2})} & 0 & 0 \\ 0 & 0 & \frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$




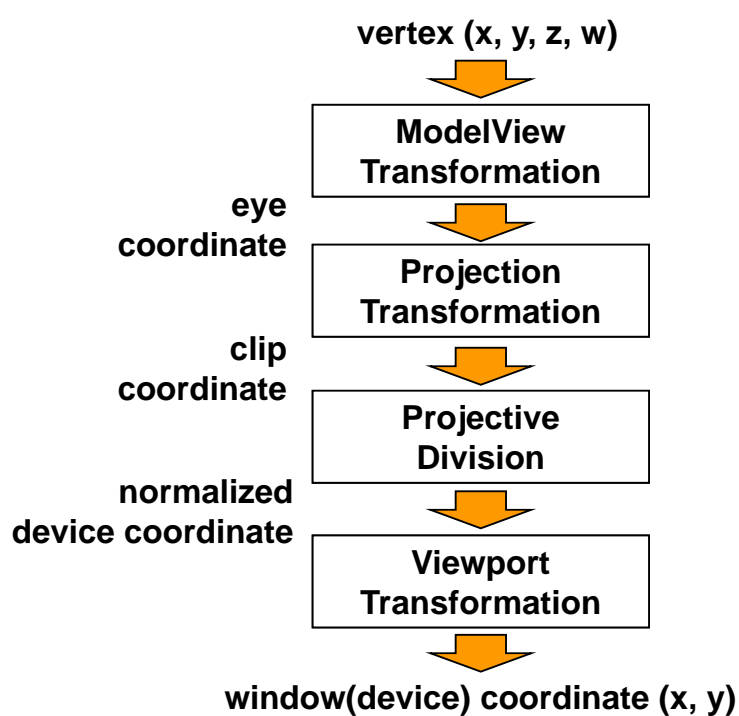
Projection Normalization





Display pipeline



Rendering pipeline (geometry stage)




Viewport



Viewport

- Usually same as window size
- Aspect ratio = width/height



`glViewport(x,y,width,height)`

Faux Plafond - Cosmic Promenade - Mikros Image

□ Siggraph 2000



□ http://www.siggraph.org/publications/video-review/SVR_2000/134/

Exercise

□ Represent the following rotation using ..

1. Matrix
2. Euler angle / Fixed angle
3. Angle and Axis
4. Unit quaternion

