

Sampling

Chapter 16

Computer Graphics

- Concerns with all aspects of producing **pictures** or **images** *using a computer*

rendering process


- This chapter deals with the ‘image’
 - Picture → collection of pixels
 - ‘.. is and artifact that depicts or records visual perception’
 - Continuous image $I(x_w, y_w)$: a bivariate function
 - Discrete image $I[i][j]$: two dimensional array of color values
 - We associate each pair of integers i, j , with the continuous image coordinates $x_w = i$ and $y_w = j$

Sampling

- Point sampling: continuous vector → discrete pixel
- Our scenes are described with triangles giving a continuous 2D color field.
- Our images are digital/discrete made up of a grid of dots.
- Need to make a bridge between these two worlds (continuous vs. discrete).
- Else we will get some unnecessary artifacts called “aliasing” artifacts.
 - Jaggies, moire patterns, flickering

Sampling

- These occur when there is too much detail to fit in one pixel.
- We can mitigate these artifacts by averaging up the colors within a pixel's square.
- This is called *anti-aliasing*.




Aliasing

- The simplest and most obvious method to go from a continuous to a discrete image is by *point sampling*.
- To obtain the value of a pixel i, j , we sample the continuous image function at a single integer valued domain location: $I[i][j] \leftarrow I(i, j)$
- This can results in unwanted artifacts.


Aliasing

- Scene made up of black and white triangles: jaggies at boundaries
 - Jaggies will crawl during motion
- If triangles are small enough then we get random values or weird patterns
 - Will flicker during motion
- The heart of the problem → too much information in one pixel



Anti-aliasing

- Intuitively: the single sample is a bad value, we would be better off setting the pixel value using some kind of average value over some appropriate region.



Anti-aliasing

- Mathematically this can be modeled using *Fourier analysis*.
 - Breaks up the data by “frequencies” and figures out what to do with the un-representable high frequencies.

- Nyquist Sampling Theorem**


- “The sampling frequency should be at least twice the highest frequency contained in the signal.”

$$f_s \geq 2f_c$$




Example


- An input signal with $f = 1$ Hz



- If we sample by more than 2 Hz, we can reconstruct the shape correctly




2 Hz




3 Hz

Example

- If we sample by 1.5 Hz (≤ 2 Hz), there might be an ambiguity about the signal shape.



1.5 Hz



Anti-aliasing

- We can also model this as an optimization problem.
- These approaches lead to:

$$I[i][j] \leftarrow \int \int_{\Omega} dx dy I(x, y) F_{i,j}(x, y)$$

- where $F_{i,j}(x, y)$ is some function that tells us how strongly the continuous image value at $[x, y]^t$ should influence the pixel value i, j
- In this setting, the function $F_{i,j}(x, y)$ is called a filter.
 - In other words, the best pixel value is determined by performing some continuous weighted averaging near the pixel's location.
 - Effectively, this is like blurring the continuous image before point sampling it.


Box Filter

- We often choose the filters $F_{i,j}(x, y)$ to be something non-optimal, but that can more easily be computed with.
- The simplest such choice is a *box filter*, where $F_{i,j}(x, y)$ is zero everywhere except over the 1-by-1 square center at $x = i, y = j$.




- Calling this square $\Omega_{i,j}$, we arrive at
$$I[i][j] \leftarrow \int \int_{\Omega_{i,j}} dx dy I(x, y)$$
- In this case, the desired pixel value is simply the average of the continuous image over the pixel's square domain.

Filters



convolution

<http://mathworld.wolfram.com/Convolution.html>



Green curve is the convolution of the blue and red curves as a function of t .






Diagram 1: The source pixel and its surrounding pixels are all mathematically merged to produce a single destination pixel. The matrix slides across the surface of the source image, producing pixels for the destination image.

convolution


identity

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$



edge detection

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



sharpen

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



box blur

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$




Over-sampling

- Even that integral $I[i][j] \leftarrow \int \int_{\Omega_{i,j}} dx dy I(x, y)$ is not really reasonable to compute.

- Instead, it is approximated by some sum of the form:

$$I[i][j] \leftarrow \frac{1}{n} \sum_{k=1}^n I(x_k, y_k)$$

a pixel




where k indexes some set of locations (x_k, y_k) called the sample locations.

- The renderer first produces a “high resolution” color and z-buffer “image”,
 - where we will use the term *sample* to refer to each of these high resolution pixels.

Over-sampling

- Then, once rasterization is complete, groups of these samples are averaged together, to create the final lower resolution image.



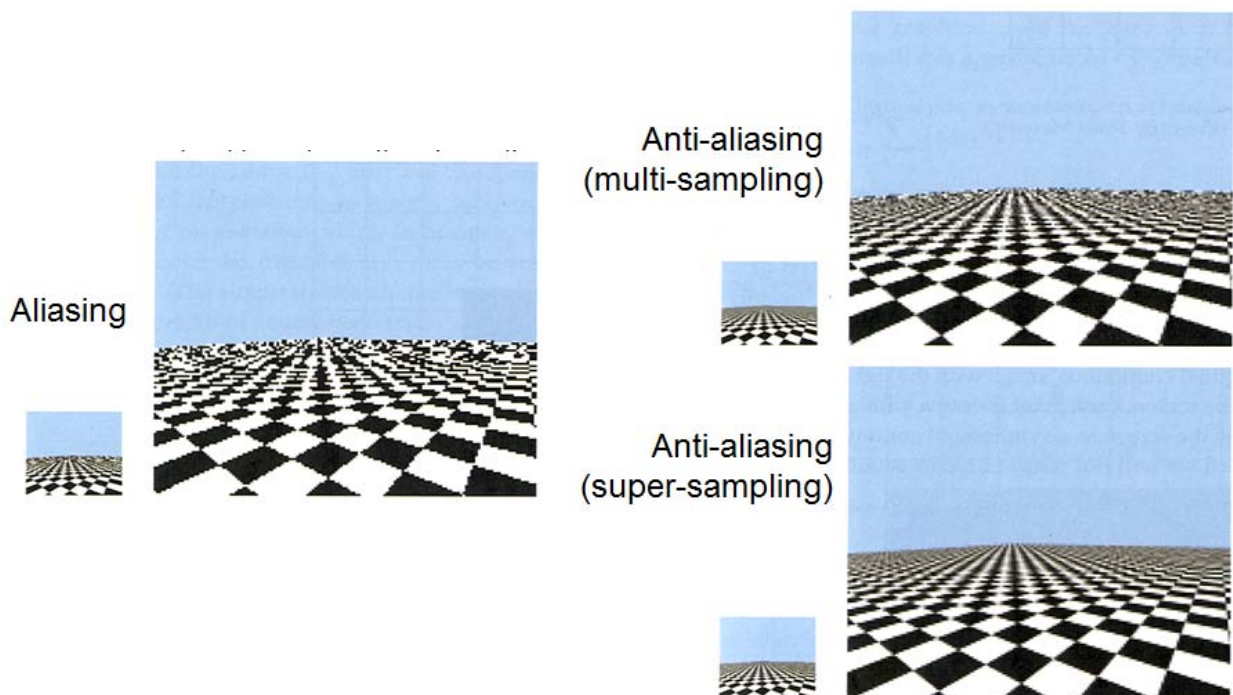
Super-sampling

- If the sample locations for the high resolution image form a regular, high resolution grid, then this is called *super sampling*.
- We can also choose other sampling patterns for the high resolution “image”,
 - Such less regular patterns can help us avoid systematic errors that can arise when using the sum to replace the integral.

Multi-sampling

- In OpenGL, we can also choose to do *multisampling*.
- Render to a “high resolution” color and z-buffer
- *During the rasterization* of each triangle, “coverage” and z-values are computed at this (high)sample level.
 - But for efficiency, the fragment shader is only called **only once per final resolution pixel**.
 - This color data is shared between all of the samples hit by the triangle in a single (final resolution) pixel.
 - Once rasterization is complete, groups of these high resolution samples are averaged together.
- Multisampling can be an effective anti-aliasing method since, without texture mapping, colors tend to vary quite slowly over each triangle, and thus they do not need to be computed at high spatial resolution.
 - → Mipmapping (next chapter)

Aliasing vs. anti-aliasing




Camera

- In digital cameras, anti-aliasing is accomplished by a combination of the spatial integration that happens over the extent of *each pixel sensor*, as well as by the *optical blurring* that happens at due to the lens.
- Some cameras also include additional optical elements specifically to blur the continuous image data before it is sampled at the sensors.

Image compositing

- Given two discrete images, a foreground, I^f , and background, I^b , that we want to combine into one image I^c .
- Simple: in composite, use foreground pixels where they are defined. Else use background pixels.
 - This will give us a jagged boundary.
- Real image would have “boundary” pixels with blended colors.
 - But this requires using “sub-pixel” information.



Alpha blending

- Associate with each pixel in each image layer, a value $\alpha[i][j]$ that describes the overall opacity or coverage of the image layer at that pixel.

- An alpha value of 1 represents a fully opaque/occupied pixel, while a value of 0 represents a fully transparent/empty one.
- A fractional value represents a partially transparent (partially occupied) pixel.

- Alpha will be used during **compositing**.

- $I(x, y)$ continuous image $I[i][j] \leftarrow \int \int_{\Omega_{i,j}} dx dy I(x, y) C(x, y)$
- $C(x, y)$ binary valued coverage $\alpha[i][j] \leftarrow \int \int_{\Omega_{i,j}} dx dy C(x, y)$
(x,y) domain (1 "occupied"

Over operation

- To compose $I^f[i][j]$ *over* $I^b[i][j]$ we compute the composite image colors $I^c[i][j]$ using


$$I^c[i][j] \leftarrow I^f[i][j] + I^b[i][j](1 - \alpha^f[i][j])$$

- the amount of observed background color at a pixel is proportional to the transparency of the foreground layer at that pixel.
- Alpha for the composite image is computed as

$$\alpha^c[i][j] \leftarrow \alpha^f[i][j] + \alpha^b[i][j](1 - \alpha^f[i][j])$$

Over operation

- If background is opaque, so the composite pixel is opaque.
- But we can model more general case as part of blending multiple layers.
 - a reasonable approximation to the correctly rendered image.



Over properties

- over operation is associative

$$I^a \text{ over } (I^b \text{ over } I^c) = (I^a \text{ over } I^b) \text{ over } I^c$$

but not commutative.

$$I^a \text{ over } I^b \neq I^b \text{ over } I^a$$

In Practice

- In OpenGL, alpha is used not just for image compositing, but in more general sense as tool for modeling transparency and for blending color values.
- (R,G, B, A) → `fragColor` in fragment shader
 - `glEnable(GL_BLEND)`
 - `glBlendFunc`

Figure 16.5:
A furry bunny is drawn as a series
of larger and larger concentric bunnies.
each is partially transparent.
Amazingly, this looks like fur.



Blending


- Combine pixels with what's in already in the framebuffer

- In OpenGL

- glEnable(GL_BLEND)

- glBlendFunc(source_factor, destination_factor)

GL_ONE, GL_ZERO, GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA,
GL_DST_ALPHA, GL_ONE_MINUS_DST_ALPHA



jinah@cs.kaist.ac.kr

CS380 (Spring 2016)

29

Compositing Techniques

- Alpha blending

Some Applications

- Multipass rendering

- Blending allows results from multiple drawing passes to be combined together

- Image compositing

- merge a set of images into a single image


- Antialiasing

- alpha value computed by computing sub-pixel coverage

- Depth Cueing and Fog

- C: color, f: fog factor

$$C_{s'} = f C_s + (1-f) C_f$$



<fog demo>

jinah@cs.kaist.ac.kr

CS380 (Spring 2016)

30

Multipass Rendering

□ Motion blur



□ Depth of field



jinah@cs.kaist.ac.kr

CS380 (Spring 2016)

Today's animation

□ PGI-13

- Siggraph 2004
- Parental Guidance for Certain Imaginations for Children under the age of 13....
A scary imagination comes from a sudden curiosity about the materials of a tea bag before put into a cup of hot water. ...
- Average CPU time for rendering per frame: 18 minutes (10 min ~ 3 hours).
- Many use of alpha-channel sequences were needed.

